

---

# MNIST CNN Accelerator Design on Zynq-7000 SoC

---

## Presenter Names:

**V. Ashish Nadh - VU22EECE0100389**

**P. Deepika - VU22EECE0100129**

**Y. Manohar - VU22EECE0100099**

**B. Aryan Ratnakar - VU22EECE0100175**

## Project Guide :

**Dr. G. V. K. Sharma**

**Department of Electrical, Electronics and Communication Engineering**

GITAM School of Technology,

GITAM (Deemed to be University), Visakhapatnam.

# Table of Contents

**Introduction & Problem Understanding**

**Methodology**

**CNN Architecture:**

- Why we choose the architecture
- Proposed 7-layer CNN Architecture
- Detailed layer configuration

**ZEDBOARD ARCHITECTURE AND EXERCISES:**

- ZedBoard Exercises-(1,2,3)
- ZedBoard Exercise-4
- Zedboard Exercise-5

**Future Work**

**References**

**Thank you**

# Introduction & Problem Understanding

**Application Context:** This project implements a hardware accelerator for a Convolutional Neural Network (CNN) specifically designed for **handwritten digit detection and recognition**. The target model is optimized to achieve high accuracy on the **MNIST dataset**, which consists of 28x28 pixel grayscale images.

**Technical Problem:** The core technical challenge is accelerating the massive computational load of a CNN – particularly the convolution and fully connected layers – within the strict power and resource constraints of an edge device. Relying on the CPU alone is too slow, while a GPU is too power-hungry.

**Project's Technical Goal:** The goal is to design a low-power, high-efficiency hardware accelerator in the FPGA's Programmable Logic (PL). This accelerator will offload the heavy computations from the ARM Cortex-A9 processor, following a **hardware-software co-design** approach on the ZEDBoard. Success is defined by minimizing power-hungry external memory accesses through on-chip data reuse.

Our proposed methodology follows a structured hardware-software co-design process, broken down into distinct phases. This approach allows for systematic development, testing, and validation of the CNN accelerator on the Zynq-7000 SoC platform.

## Phase 1: CNN Model Design & Architecture :

- Design of Convolution Neural network CNN in detail ( Convolution layer, pooling layer, Fully connected layers, Activation Functions etc. )

## Phase 2: RTL Design & Functional Simulation

- **Modular Hardware Design:** Develop individual Verilog modules for each core function: Convolution Layer, Activation Function , Pooling Layer, Fully Connected Layer and a top-level controller.
- **Behavioral Simulation:** Create a Verilog testbench to simulate and verify the functional correctness of each module independently and then the integrated accelerator as a whole.

# Methodology

## Phase 3: System Integration with Zynq SoC

- **IP Core Creation:** Package the verified RTL design into an IP (Intellectual Property) core using Xilinx Vivado.
- **Block Design:** Integrate the custom CNN accelerator IP with the Zynq-7000's ARM Processing System (PS) using the standard AXI bus for communication.

## Phase 4: Software Development & On-Board Validation

- **Software Development :** Develop a C/C++ application in the Xilinx SDK to run on the ARM processor. This application will manage data transfer to/from the accelerator and control its operation.
- **System Testing:** Generate the final bitstream, program the FPGA, and run the complete system to perform inference on MNIST test images.
- **Verification:** Validate the hardware's classification output against the results from a software model to confirm correctness.

# **CNN ARCHITECTURE**

## **7-Layer Modified LeNet-5 Architecture**

# Why We Chose This Architecture?

- **Hardware Constraints:** The ZedBoard (Zynq-7000) has limited On-Chip Memory (BRAM) and DSP slices.
- **Bottleneck Analysis:** Standard models like VGG-16 (138M parameters) are too heavy and require slow external DDR memory access.
- **Our Solution:** We selected a customized 7-Layer Architecture (Modified LeNet-5).
- **Key Benefit:** Extremely lightweight (~60k parameters), allowing the entire model to fit inside the FPGA's fast internal BRAM for real-time processing.

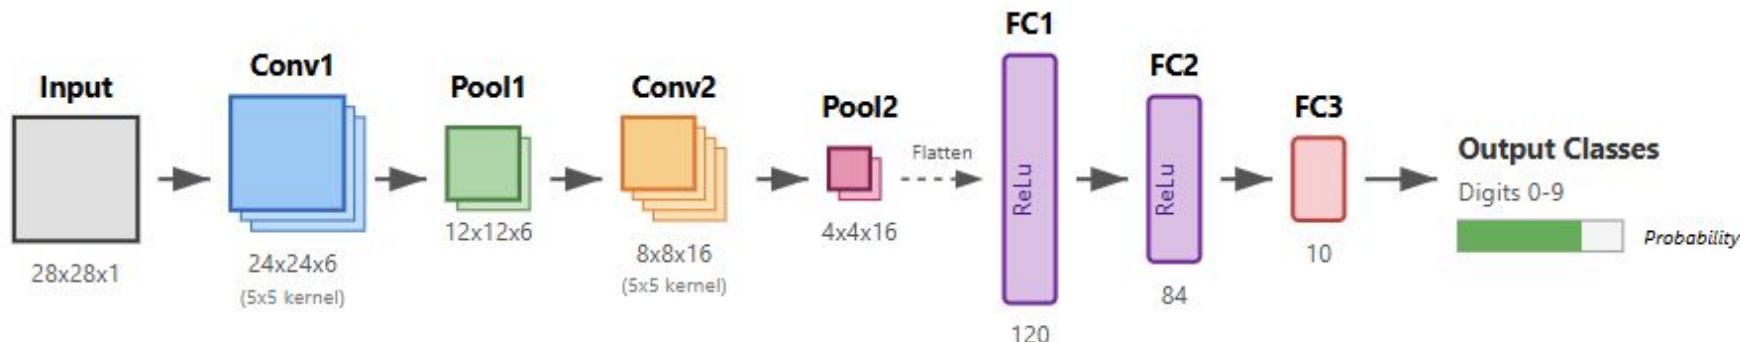
# Proposed 7-Layer CNN Architecture

- **Input:** 28 X 28 Grayscale Image.
- **Feature Extraction:** 2 Convolutional Layers (to detect edges and shapes).
- **Downsampling:** 2 Max-Pooling Layers (to reduce data size by 75% at each step).
- **Classification:** 3 Fully Connected Layers (mapping features to the 10 digits).
- **Final Output:** Probability vector for Digits 0-9.

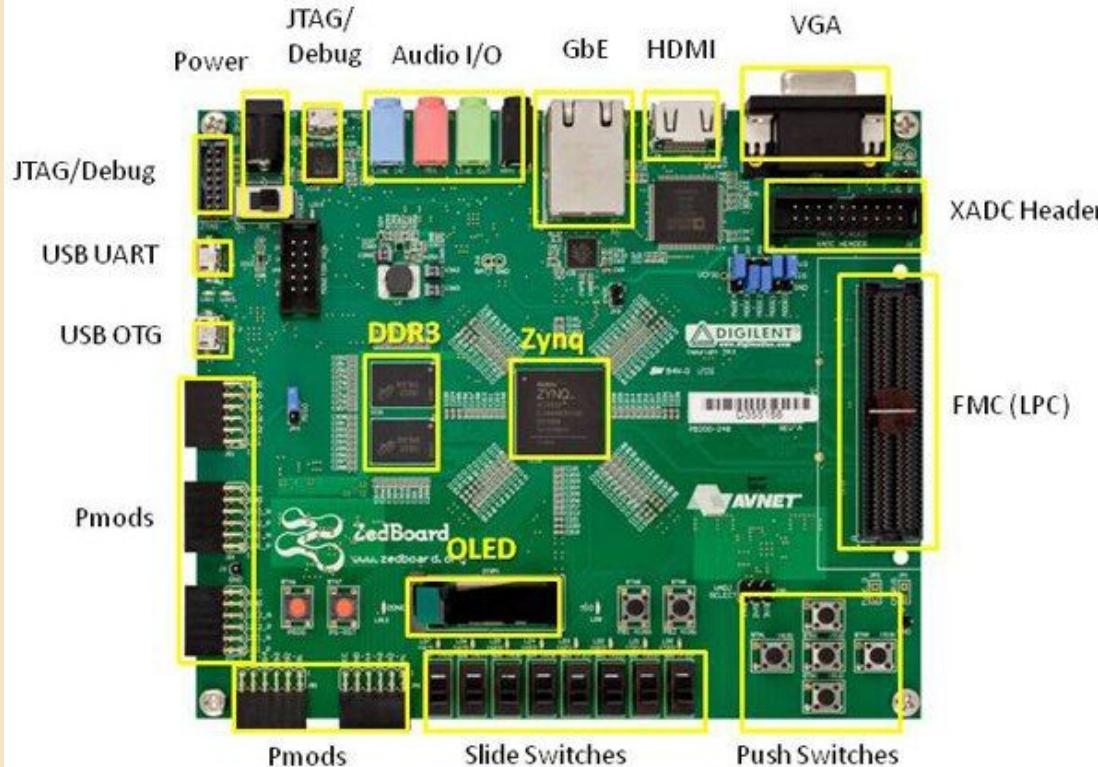
# Detailed Layer Configuration

- **Layer 1 (Conv1):** 6 Filters, Kernel Size 5 X 5, Output: 24 X 24 X 6.
- **Layer 2 (Pool1):** Max Pooling, Window 2 X 2 (Stride 2), Output: 12 X 12 X 6.
- **Layer 3 (Conv2):** 16 Filters, Kernel Size 5 X 5, Output: 8 X8 X 16.
- **Layer 4 (Pool2):** Max Pooling, Window 2 X 2 (Stride 2), Output: 4 X 4 X 16.
- **Fully Connected Layers:**
  - **FC1:** 120 Neurons
  - **FC2:** 84 Neurons
  - **Output:** 10 Neurons (Softmax equivalent)

# 7-Layer CNN Architecture for ZedBoard (Modified LeNet-5)



# ZEDBOARD ARCHITECTURE AND EXERCISES



\* SD card cage and QSPI Flash reside on backside of board

## ZEDBOARD ARCHITECTURE

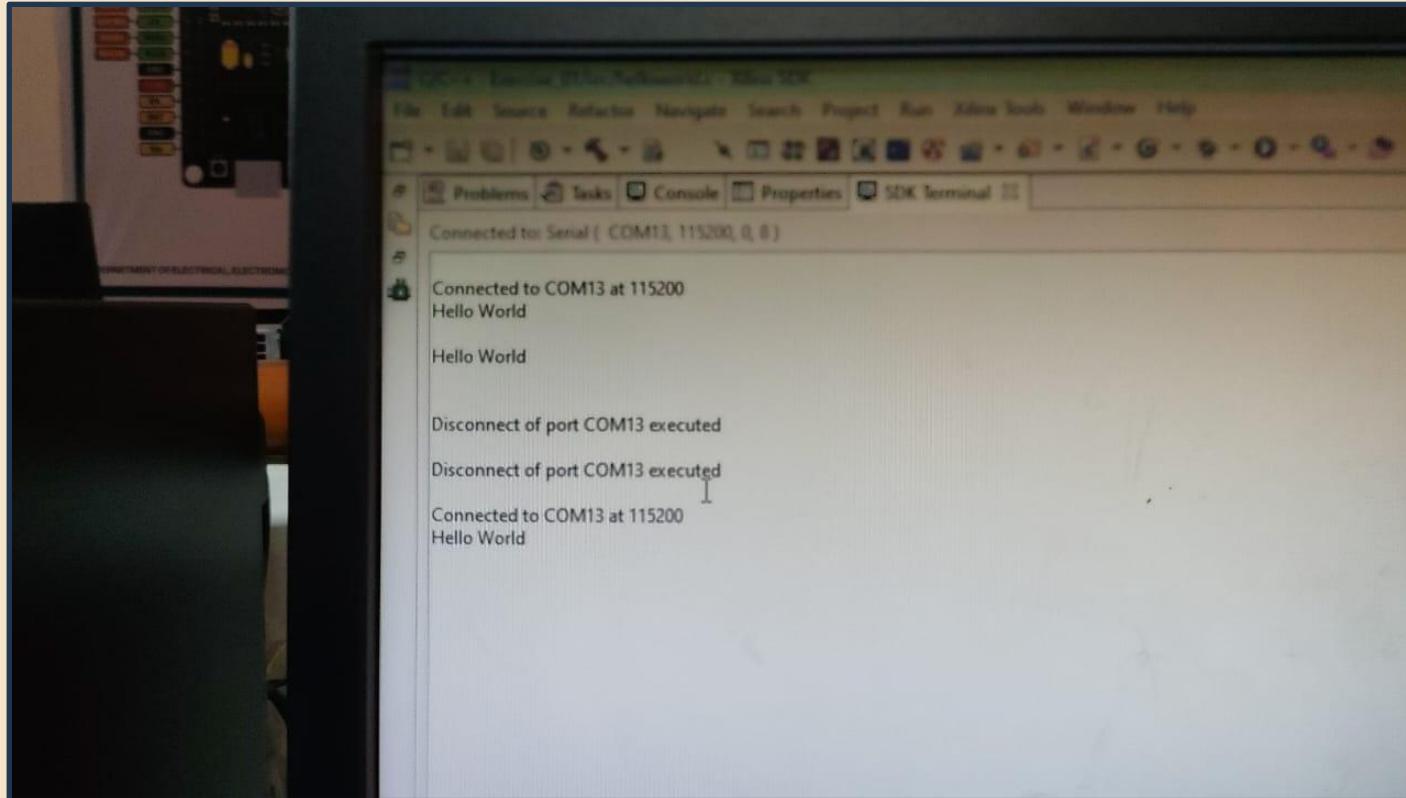
# ZED BOARD Exercises (1,2,3)

**Objective:** To verify the toolchain (Vivado/SDK) and establish control over the ARM Cortex-A9 Processing System.

## Implementation:

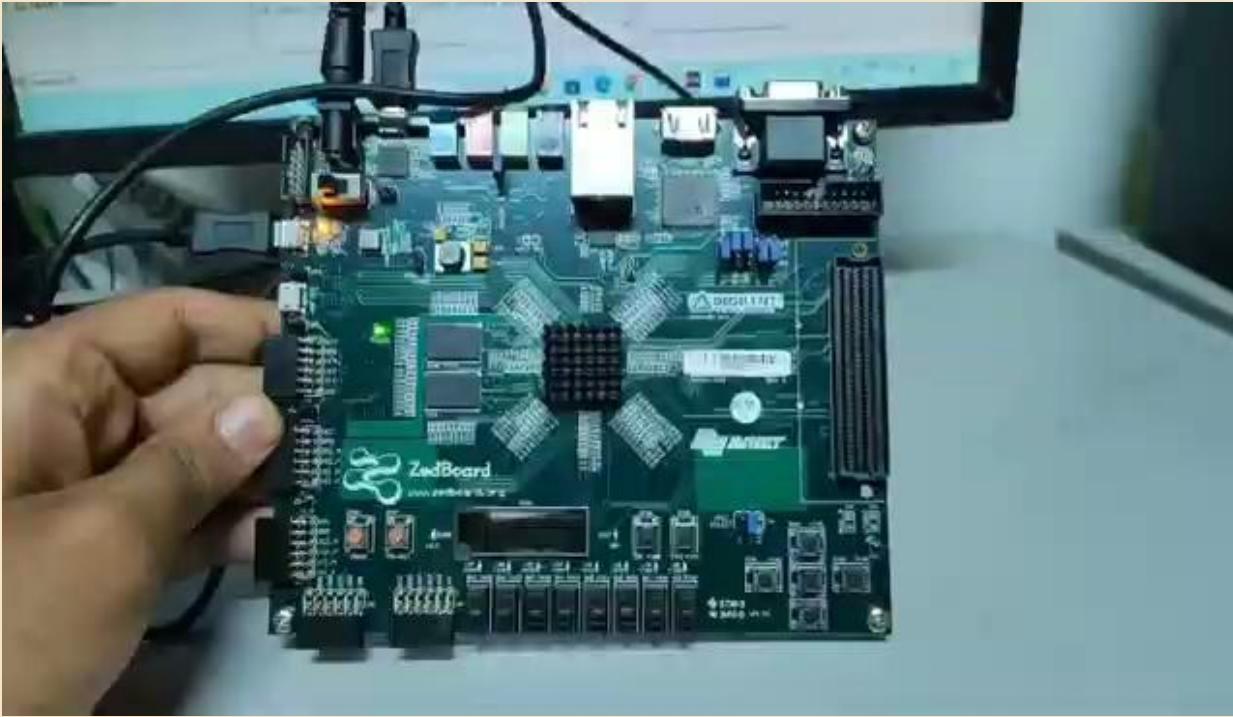
- **Hardware:** Configured the **ZYNQ7 Processing System** with ZedBoard presets (DDR Memory & Clock generation).
- **Software (BSP):** Utilized Xilinx Board Support Package (BSP) drivers.
- **Debugging environment:** Explored the Debugging environment in Xilinx SDK
- **Key Result (Exercise 1 & 2):**
  - verified the functionality of the Zynq chip's Processing System (PS) on the ZedBoard through "Hello World" in the SDK terminal.
  - Bypassed the FPGA fabric to control **MIO Pin 7** directly from the Processor.
  - Successfully toggled User LED **LD9**.

# ZEDBOARD Exercise 1



SDK Terminal showing "Hello World"

# ZEDBOARD Exercise 2



LED (LD9) Blinking on the Processing System (PS)

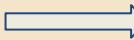
# ZEDBOARD Exercise 4

## Hardware-Software Architecture

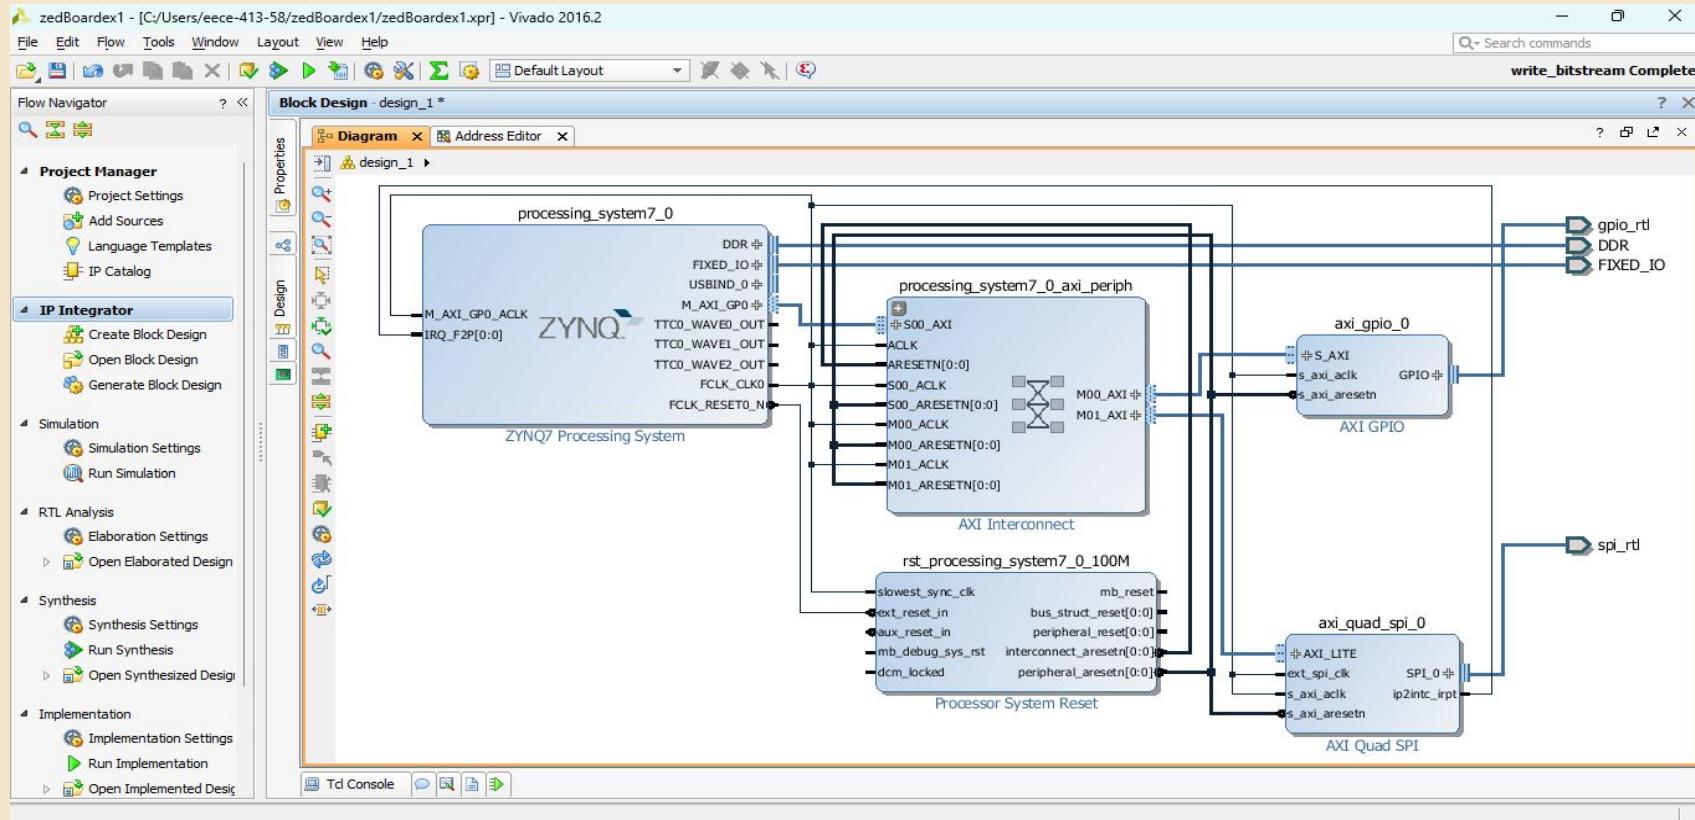
### Objective:

To establish a communication pipeline between the ARM Cortex-A9 Processor (PS) and Custom FPGA Logic (PL).

### System Architecture:

- **The Bridge (AXI Interface):**
  - Enabled **M\_AXI\_GP0** (General Purpose Master) port on the Zynq.
  - This converts software commands into hardware signals.
- **The Hardware (AXI GPIO IP):**
  - Integrated a **16-bit AXI GPIO** block into the PL fabric.
  - Configured as a **Bidirectional Bus** to handle both input and output data simultaneously.
- **Physical Mapping (Constraints):**
  - **Input Map:** Bits [0-7]  Physical DIP Switches (F22 - M15).
  - **Output Map:** Bits [8-15]  Physical LEDs (T22 - U14).

# Vivado IP BLOCK Design

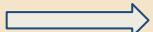


# ZEDBOARD Exercise 5

## Interactive System Logic & Results

Control Algorithm (The "Loop"):

The system operates on a continuous feedback loop:

1. **INPUT (Hardware  Software):**
  - CPU reads the state of the **8 DIP Switches** from the PL addresses.
2. **PROCESS (Software Logic):**
  - CPU performs a **Left Bit-Shift ( $<< 8$ )** operation.
  - This moves the data from the "Input" bit positions (0-7) to the "Output" bit positions (8-15).
3. **OUTPUT (Software  Hardware):**
  - CPU writes the processed value back to the **LED Registers**, instantly lighting the corresponding LEDs.

## Experimental Outcome:

- **Latency:** Real-time response (Zero perceptible delay).
- **Verification:** Verified that **SW[0] ON  LED[0] ON**, confirming successful bidirectional data transfer.

# ZEDBOARD Exercise 5



LEDs Blinking by flipping DIP Switches on the Programmable logic (PL)

## Memory Interfacing & Direct Register Access

**Goal:** To understand the memory-mapped architecture of the Zynq-7000 SoC and perform direct I/O operations without high-level drivers.

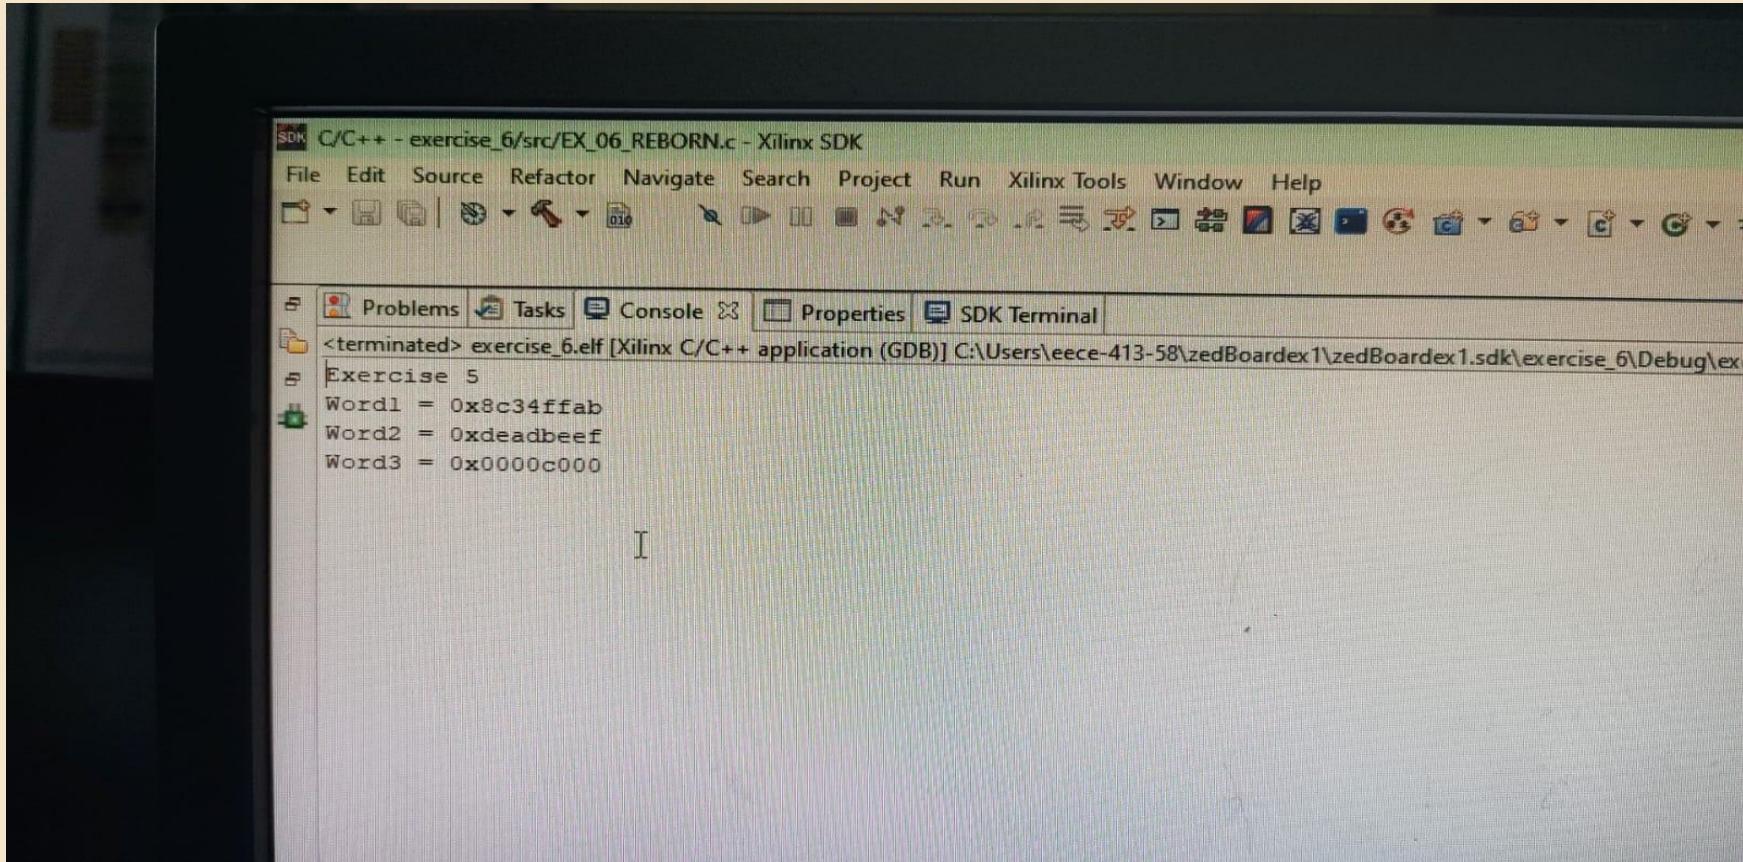
### Technical Details

- **Memory-Mapped I/O:** Every peripheral and memory block (OCM, DDR, GPIO) is assigned a unique physical address. We used `xil_io.h` for low-level access.
- **On-Chip Memory (OCM):** Explored the high-speed 256KB RAM inside the Processing System (PS).
- **Direct Address Manipulation:** We targeted address `0xE000A064` to read the raw state of the PS-side GPIO pins.

### Experimental Results (Output)

- **Word 1:** `0x8C34FFAB` - Successfully wrote and verified a 32-bit pattern in OCM.
- **Word 2:** `0xDEADBEEF` - Verified successful word-aligned memory transactions.
- **Word 3 (Interaction):** `0x0000C000` - When **Push Buttons (PB1/PB2)** were pressed, we observed real-time changes in bits 14 and 15 of the register.

# ZEDBOARD Exercise 6



The screenshot shows the Xilinx SDK C/C++ IDE interface. The title bar reads "SDK C/C++ - exercise\_6/src/EX\_06\_REBORN.c - Xilinx SDK". The menu bar includes File, Edit, Source, Refactor, Navigate, Search, Project, Run, Xilinx Tools, Window, and Help. Below the menu is a toolbar with various icons. The central workspace has tabs for Problems, Tasks, Console, Properties, and SDK Terminal. The Console tab is active, displaying the output of a GDB session for "Exercise 5". The output shows variable definitions:

```
<terminated> exercise_6.elf [Xilinx C/C++ application (GDB)] C:\Users\eece-413-58\zedBoardex1\zedBoardex1.sdk\exercise_6\Debug\exe
Exercise 5
Word1 = 0x8c34ffab
Word2 = 0xdeadbeef
Word3 = 0x0000c000
```

A cursor is visible in the console window.

## Hardware Timers (Polled Mode)

**Goal:** To achieve precise time measurement using the Private SCU (Snoop Control Unit) Timer instead of software delay loops.

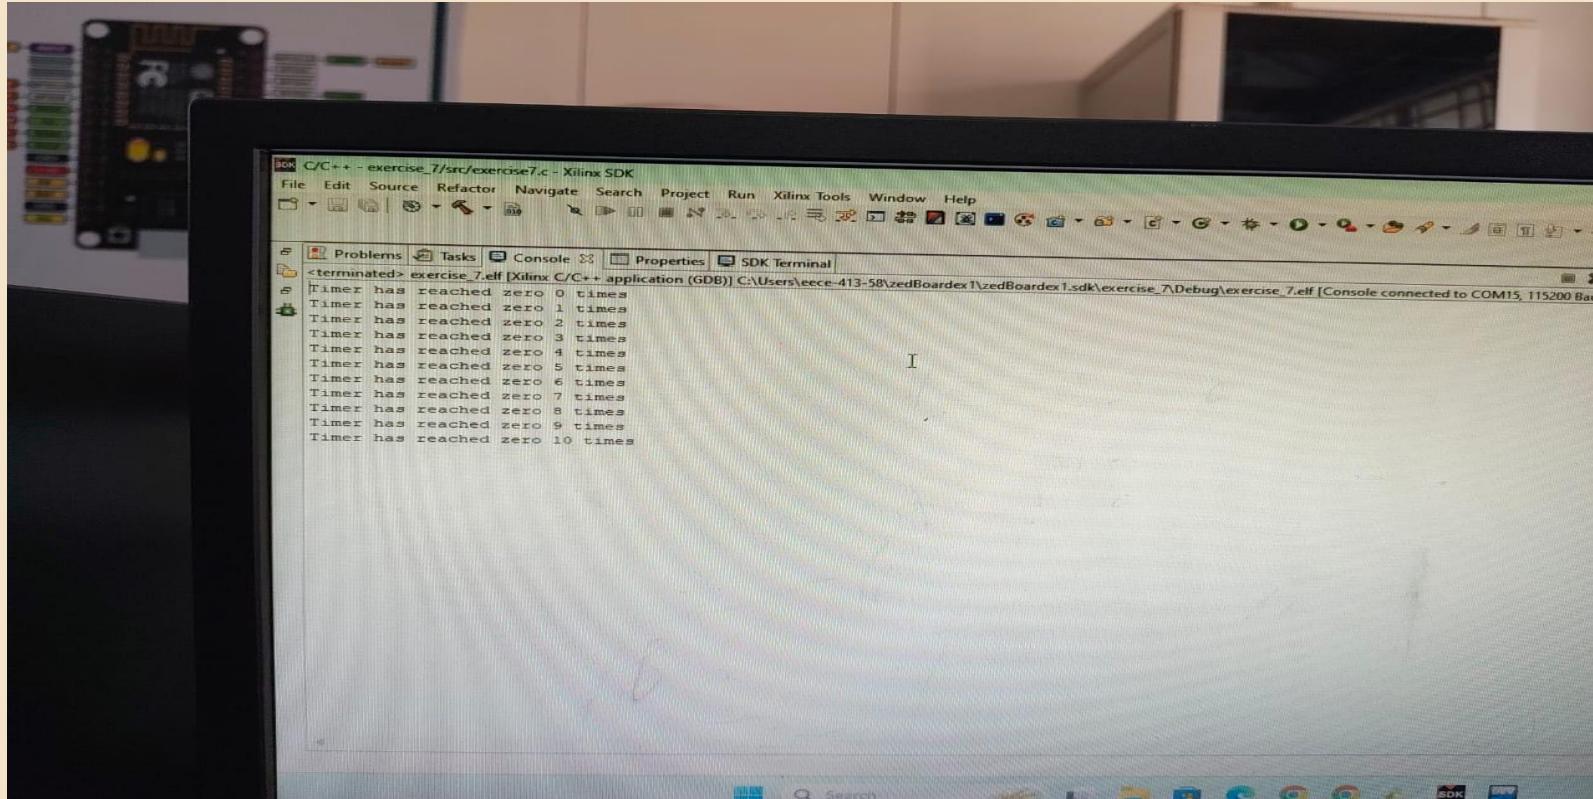
### Technical Details

- **SCU Timer:** A 32-bit hardware counter that decrements at half the CPU clock frequency (333 MHz on the ZedBoard).
- **Polled Mode Logic:** The CPU enters a loop that continuously checks the "Timer Counter" register. Once it reaches zero, the event is triggered.
- **Calibration:** Used `XScuTimer_Config` to initialize the hardware and set a precise 1-second (1Hz) interval.

### Experimental Results (Output)

- **Terminal Display:** "the timer has reached zero 0 times" ... up to "the timer has reached zero 10 times."
- **Precision:** Unlike software delays, this output remained perfectly consistent at exactly 1000 ms intervals.

# ZEDBOARD Exercise 7



## Real-Time Interrupt Handling

**Goal:** To implement an efficient system where the hardware "notifies" the CPU of events, allowing for true multitasking.

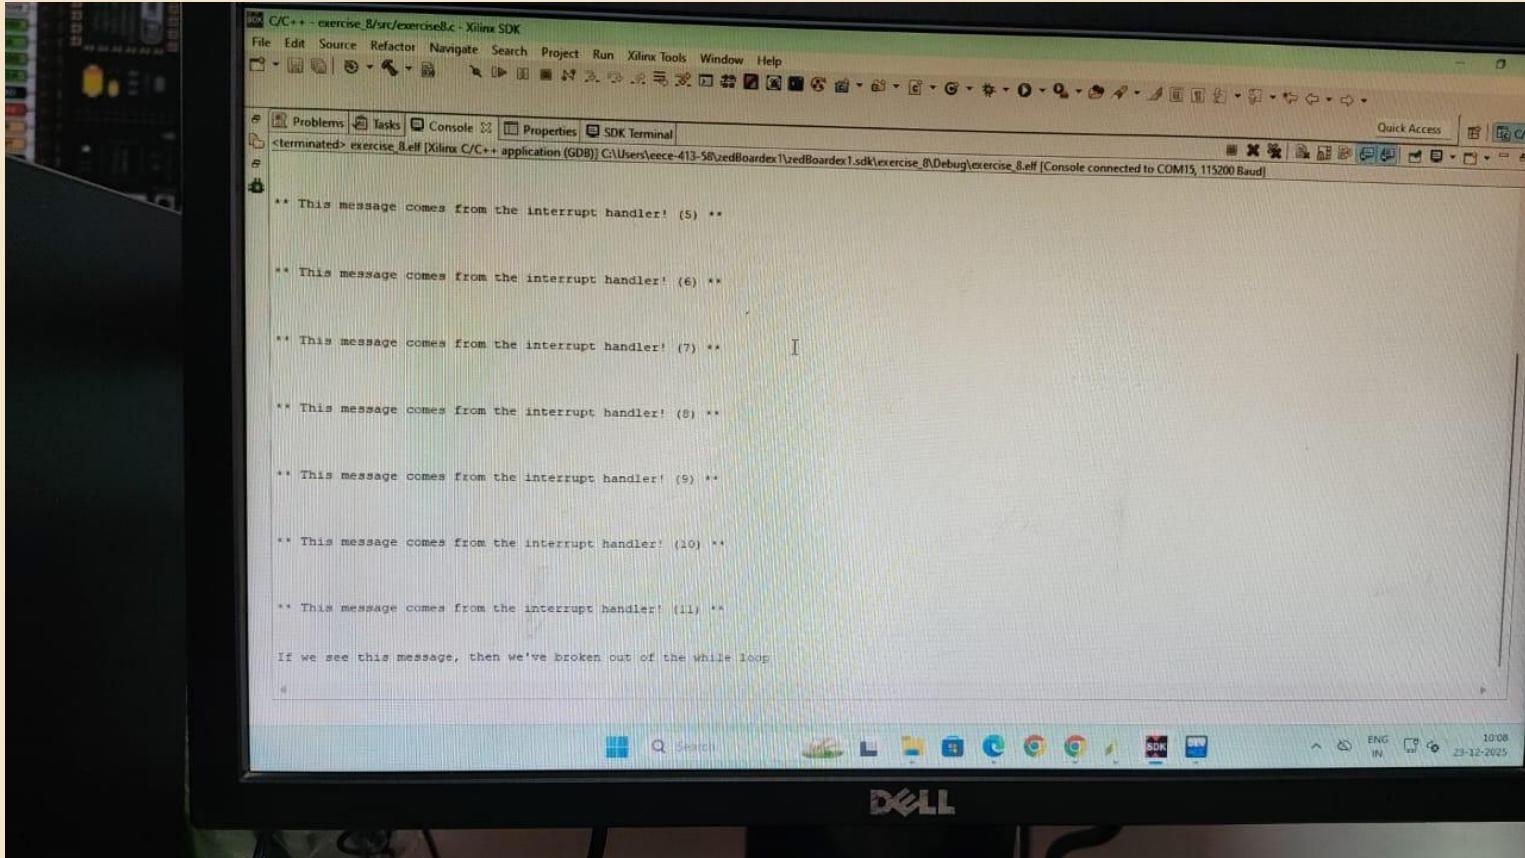
### Technical Details

- **Interrupts vs. Polling:** In polling (Ex 7), the CPU is stuck in a loop. With interrupts, the CPU is free to do other work until the hardware signals it.
- **GIC (General Interrupt Controller):** Configured the GIC to route signals from the SCU Timer to the ARM core.
- **ISR (Interrupt Service Routine):** Developed a handler function (`my_timer_interrupt_handler`) that executes only when the timer expires.

### Experimental Results (Output)

- **Multitasking Success:** 1. **Task A (Background):** Printed "message comes from interrupt handler" to the UART terminal at 1Hz. 2. **Task B (Foreground):** Continuously updated the **8 LEDs** in real-time based on the position of the **DIP Switches**.

# ZEDBOARD Exercise 8



```
** This message comes from the interrupt handler! (5) **

** This message comes from the interrupt handler! (6) **

** This message comes from the interrupt handler! (7) **

** This message comes from the interrupt handler! (8) **

** This message comes from the interrupt handler! (9) **

** This message comes from the interrupt handler! (10) **

** This message comes from the interrupt handler! (11) **

If we see this message, then we've broken out of the while loop
```

# Future Work

Our plan for completing this project is a direct execution of our methodology. The following steps outline our roadmap leading up to the final project delivery.

- **CNN Model architecture and Design - completed**
- **RTL Implementation & Simulation**
  - **Goal:** To translate the CNN architecture into hardware and verify its logic.
  - **Tasks:**
    - Develop modular Verilog code for the convolution engine, pooling/activation units, and the fully connected layer.
    - Write a comprehensive testbench to simulate the data flow through the entire accelerator.
    - Debug and validate the RTL design to ensure it is functionally correct.

# Future Work

## System Integration

- **Goal:** To integrate the verified hardware accelerator with the Zynq processing system.
- **Tasks:**
  - Package the Verilog design into an AXI4-compliant IP core in Vivado.
  - Create the complete SoC block design, connecting our IP to the ARM processor.
  - Synthesize and implement the design, ensuring it meets timing and resource constraints.

# Future Work

## Software Development & On-Board Validation

- **Goal:** To control the accelerator with software and demonstrate the complete, working system.
- **Tasks:**
  - Write a C/C++ application to send image data and control signals to the accelerator.
  - Deploy the bitstream and software application to the Zynq board.
  - Perform end-to-end testing with MNIST images and verify the classification results on a terminal.

# References

1. Jiang, J., Zhou, Y., Gong, Y., Yuan, H., & Liu, S. (2025). *FPGA-based Acceleration for Convolutional Neural Networks: A Comprehensive Review*. arXiv preprint arXiv:2505.13461.
2. Srilakshmi, S., & Madhumati, G. L. (2023). *A Comparative Analysis of HDL and HLS for Developing CNN Accelerators*. 2023 International Conference on Intelligent Systems and various other Computer Applications (ICISVCA).
3. García-López, J., et al. (2020). *A hardware accelerator for the inference of a Convolutional Neural Network*. Ingeniería, 25(1), 107-118.
4. Khan, M. A., et al. (2022). *A Step-by-Step Re-Configurable CNN Engine Design on FPGA*. Electronics, 11(23), 3883.

---

**Thank  
You**